Mela Task Radar

Comprehensive Project Review & Analysis

Architecture · Logic · Issues · Recommendations

Report Type:	Full codebase review - logic, issues, improvements
Date:	May 04, 2026
Scope:	Backend (FastAPI/Python) · Frontend (Next.js/TS) · Infra
Status:	MVP-level - pre-production credential validation required
Reviewer:	Claude AI Code Analysis

1. Executive Summary

Mela Task Radar is an AI-powered task-intelligence platform for Microsoft 365. It automatically discovers implicit work items buried in Outlook emails and Teams messages, extracts structured task metadata via Azure OpenAI GPT-5.2, stores results in a relational database, and propagates them to Excel workbooks and Microsoft Planner. A REST API and an MCP (Model Context Protocol) server expose all capabilities to external AI agents such as Mela AI.

This review covers every layer of the system: authentication, data ingestion, AI extraction, task persistence, sync targets, the Next.js frontend, infrastructure configuration, test coverage, and production readiness.

Overall Assessment

Architecture	Clean, async-first, well-layered. Strong separation of concerns.
Code Quality	Above-average for an MVP. Type hints, audit logs, request context throughout.
Test Coverage	37+ unit & integration tests. Mocked Graph and OpenAI. Real-credential gaps remain.
Security Posture	OAuth 2.0 + Fernet token encryption + httpOnly JWT cookie. No rate limiting.
Production Readiness	All 28 MVP success criteria implemented. Requires Azure AD app + real API validation.
Top Risk	In-process worker (dev mode) silently fails if queue worker dies; not production-safe.
Teams Scanning	Stub only - disabled by feature flag. Phase 2.
Realtime Webhooks	Stub only - zero implementation. Phase 2.

Score Card

Core Scan Pipeline	9/10 - Robust with dedup, retry, audit.
AI Extraction	8/10 - Structured output with repair retry; prompt tuning needed.
Auth & Token Management	8/10 - Correct MSAL flow; PKCE could be stronger.
Excel Sync	7/10 - Works; chunked; no conflict resolution.
Planner Sync	7/10 - Approval-first pattern is safe; needs error surfacing.
Frontend UX	6/10 - Functional but thin. No loading skeletons, limited feedback.
MCP Server	7/10 - 9 tools, multi-user caveat in single-user auto-resolution.
Test Suite	7/10 - Good coverage of happy paths; edge cases sparse.
Observability	5/10 - Audit log + App Insights hook; no metrics, no tracing spans.
Documentation	9/10 - Comprehensive docs/, readiness report, architecture diagram.

2. Technology Stack & Dependencies

Full inventory of frameworks, libraries, and cloud services

Backend (apps/api)

Runtime	Python 3.12, Uvicorn 0.30.6 (ASGI)
Framework	FastAPI 0.115.6 - async request handling, OpenAPI docs auto-generated
ORM / DB	SQLAlchemy 2.0.35 async + Alembic 1.13.3 migrations
DB backends	aiosqlite (dev/CI), aioodbc -> MSSQL (production)
Auth	MSAL 1.31.0, authlib 1.3.2, python-jose (JWT signing/validation)
Encryption	cryptography 43.0.1 - Fernet symmetric encryption for stored tokens
AI	openai 1.51.0 -> Azure OpenAI GPT-5.2, structured JSON extraction
HTTP Client	httpx 0.27.2 - async calls to Microsoft Graph API
Scheduler	APScheduler 3.10.4 - daily scan trigger per user timezone
Queue	In-memory (dev) / azure-servicebus (prod) via pluggable adapter
Storage	Local ./storage (dev) / azure-storage-blob (prod) via pluggable adapter
Retry	tenacity 9.0.0 - exponential backoff on Graph 429/5xx
HTML parsing	beautifulsoup4 4.12.3 + lxml 5.3.0 - email body cleanup
MCP	mcp 1.0.0 - HTTP + stdio Model Context Protocol server
Key Vault	azure-keyvault-secrets - prod secrets management
Testing	pytest 8.3.3, pytest-asyncio 0.24.0, respx 0.21.1, freezegun 1.5.1

Frontend (apps/web)

Framework	Next.js 14.2.13 - App Router, TypeScript, server components
Styling	Tailwind CSS 3.4.13 + PostCSS
Data Fetch	SWR 2.2.5 - stale-while-revalidate hooks for live data
Icons	Lucide React 0.451.0
Utilities	clsx 2.1.1 - conditional classname merging
TypeScript	5.6.2
Linting	ESLint 8.57.0 (Next.js config)

Infrastructure

Containers	Docker + Docker Compose (5 services: db, api, worker, scheduler, web)
Cloud	Azure - App Service / Container Apps, MSSQL, Blob Storage, Service Bus, Key Vault
IaC	Azure Bicep templates (infra/azure/)
CI/CD	Not yet configured (no GitHub Actions / Azure Pipelines found)
Observability	Application Insights connection string (optional env var)

Notable Version Concerns

- GPT-5.2: Azure OpenAI model ID 'gpt-5.2-chat' is not a publicly released model as of May 2026. This may be a private preview deployment name - confirm this matches your actual Azure deployment.
- MSSQL + aioodbc in production requires the ODBC 17/18 driver installed in the container image. The Dockerfile does not install it - production builds will fail without it.
- APScheduler 3.10.4 is the 3.x branch (not 4.x). 4.x has breaking API changes; staying on 3.x is fine but it should be documented.
- mcp 1.0.0 - MCP protocol is evolving rapidly. Pin and monitor for breaking spec changes.

3. Architecture & Component Design

How the system is structured and how components communicate

High-Level Data Flow

DATA FLOW

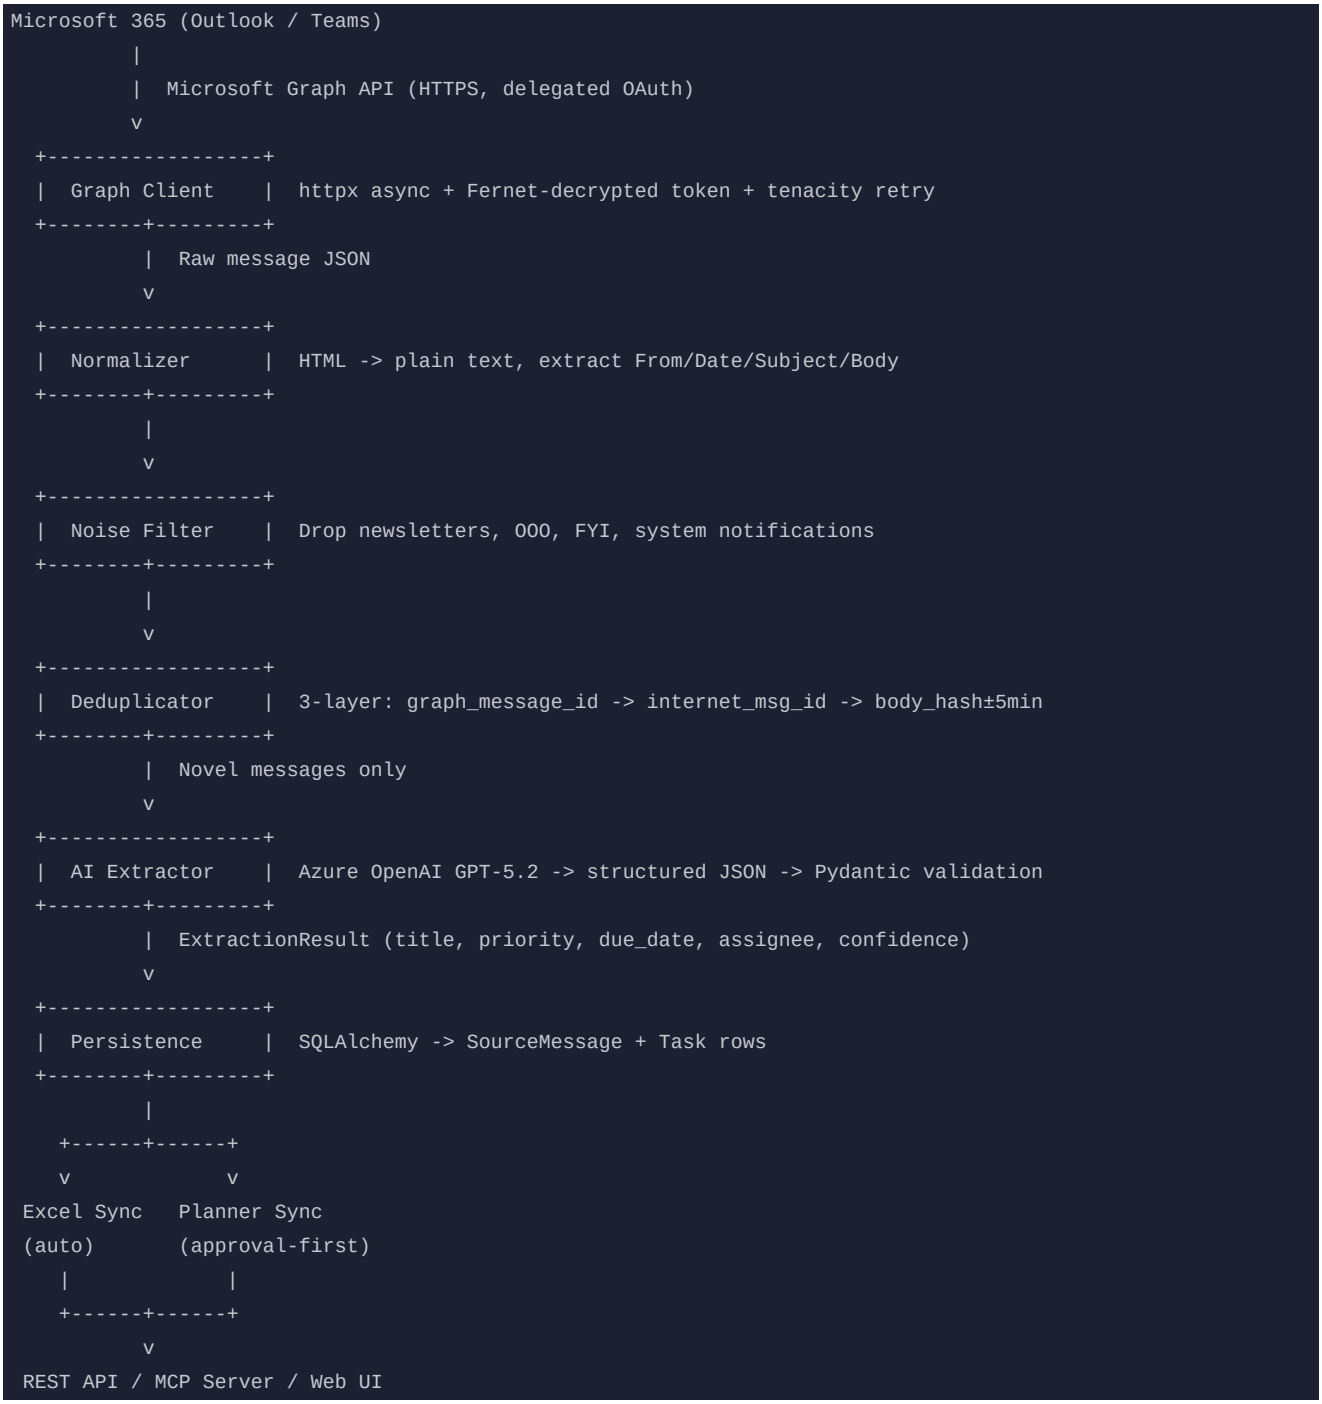

Request Context & Tenancy

Every HTTP request is enriched by deps.py which decodes the JWT session cookie and injects a RequestContext(tenant_id, user_id) into every router. Every SQLAlchemy query filters by both tenant_id AND user_id. This guarantees strict data isolation in a multi-tenant deployment backed by a single database schema - Entra tenant = one row in 'tenants'; all child tables cascade from it.

Service Layer Boundaries

routers/	HTTP layer only - validate input, call services, return responses

services/auth/	MSAL OAuth, token encryption/decryption, session JWT issuance
services/graph/	All Microsoft Graph HTTP calls - emails, teams, excel, planner
services/ai/	Azure OpenAI wrapper - prompt building, extraction, repair retry
services/tasks/	Scan orchestration, dedup, normalize, persist, audit
services/excel/	Excel workbook create/append/sync logic
services/planner/	Planner create-task with approval gate
services/queue/	Pluggable queue adapter (memory | servicebus)
services/storage/	Pluggable file adapter (local | blob)
mcp/server.py	MCP HTTP + stdio - exposes 9 tools, shares DB + services
scheduler/	APScheduler - fires daily scan jobs per user timezone

Queue Architecture

The queue is a critical seam. In development, an asyncio.Queue lives in-process alongside the FastAPI app (launched in the lifespan hook). In production, Azure Service Bus decouples the API from the worker. This pluggability is architecturally sound, but the dev mode carries a significant risk: if an unhandled exception kills the internal worker coroutine, there is no restart mechanism - scans silently stop processing while the API continues accepting requests.

4. Logic Analysis - Scan Pipeline Deep Dive

Step-by-step analysis of the core business logic

4.1 Authentication & Token Lifecycle

MSAL confidential-client OAuth 2.0 with authorization-code flow is the correct choice for delegated user consent. The implementation stores the token response encrypted with Fernet (AES-128-CBC + HMAC-SHA256) in the graph_connections table.

Token refresh logic in services/graph/client.py checks expiry - 60 seconds before every Graph call. This is correct. However, the comparison was previously broken (timezone-naive vs timezone-aware datetime) and was recently patched. The fix is correct: both sides are now forced to UTC-aware via .replace(tzinfo=timezone.utc).

4.2 Email Scanning (Outlook)

Graph endpoint: GET /me/messages with $filter and $top=50 paging. A 'high-water mark' (last_scanned_at timestamp) limits each scan to new messages only. Delta tokens (Graph Change Tracking) are NOT used - this means every scan re-fetches the filter window from Exchange, which is less efficient but simpler. For mailboxes with high volume this will cause latency creep.

- GOOD: Per-message try/except prevents one bad message from aborting the entire scan.
- GOOD: HTML body stripped via BeautifulSoup before AI processing (prevents prompt injection from HTML markup).
- CONCERN: $filter on receivedDateTime uses string comparison - timezone edge cases on DST boundaries may miss or duplicate messages.
- CONCERN: No attachment size cap before download - a large attachment could exhaust memory or disk.
- MISSING: Delta token / Graph Change Tracking - should be Phase 2 for efficiency.

4.3 Noise Filtering

services/tasks/noise.py applies keyword and sender-pattern matching to classify messages as noise (newsletters, OOO, marketing, system notifications). The filter is purely rule-based - no ML. It runs BEFORE AI extraction, saving OpenAI tokens.

- GOOD: Runs before AI - avoids wasting tokens on junk.
- GOOD: Case-insensitive matching on subject + sender.
- CONCERN: Rule-based lists will have false positives/negatives in practice. No feedback loop to improve rules.
- CONCERN: No configurable per-user noise rules - a user cannot whitelist a sender pattern.
- IMPROVEMENT: Add a 'mark as not noise' feedback endpoint so users can correct misclassifications.

4.4 Deduplication

Three-layer approach: (1) graph_message_id exact match - fastest, catches re-scans. (2) internet_message_id match - catches forward/reply chains. (3) body_hash with ±5 minute window - catches near-duplicates (same email delivered to multiple group inboxes).

- GOOD: Progressive dedup layers - fast path first.
- GOOD: Body hash window prevents duplicate tasks from mailing list replicas.
- CONCERN: Body hash is a simple hash of the cleaned body text. Two emails with identical body but different subjects (e.g., re-used template) will incorrectly dedup.
- CONCERN: The ±5 minute window is hardcoded. High-frequency senders could still produce duplicates.
- MISSING: No semantic dedup (embedding similarity) - planned for later phases.

4.5 AI Extraction (GPT-5.2)

The extractor sends a structured prompt to Azure OpenAI with response_format=json_object. The prompt instructs the model to extract: task title, description, task type, priority (high/medium/low), due date, assigned person, and confidence score (0.0-1.0).

A single repair retry is attempted on malformed JSON before failing. Confidence < 0.65 routes the task to 'needs_review' status rather than discarding it.

- GOOD: Structured JSON output format enforced at API level (not just prompt).
- GOOD: Pydantic validation ensures type safety on the extracted result.
- GOOD: Confidence threshold preserves low-confidence extractions for human review.
- GOOD: Single repair retry on malformed JSON - avoids silent failure.
- CONCERN: The system prompt content is not visible in this review - if it includes PII examples or customer data, that would be a compliance risk.
- CONCERN: No prompt versioning or A/B testing mechanism. Changes to the prompt are silent.
- CONCERN: Token cost is untracked - high-volume mailboxes could incur large OpenAI bills.
- CONCERN: 'gpt-5.2-chat' is an unconfirmed model identifier - verify against Azure deployment.
- MISSING: No caching of AI results. Re-scanning the same message (e.g., after a bug fix) will re-charge tokens unnecessarily.
- IMPROVEMENT: Log prompt token counts per scan run for cost monitoring.

4.6 Task Persistence

Successfully extracted tasks are written to the 'tasks' table alongside a 'source_messages' record. Status defaults to 'open' (or 'needs_review' for low-confidence). Every write is wrapped in an audit log entry. Sync status is tracked per-task per-target in 'task_syncs'.

- GOOD: Separate source_messages table preserves the raw message independent of the task lifecycle.
- GOOD: task_syncs table provides per-target sync audit trail.
- CONCERN: No optimistic locking - concurrent API calls (user manually triggers scan while scheduled scan is running) could create duplicate task records if dedup hasn't committed yet.
- CONCERN: Soft-delete via status='ignored' is fine for UX, but records accumulate forever. No archival/purge policy.

5. Bugs, Issues & Broken Items

Ranked by severity - critical items must be fixed before production

5.1 Critical Issues

[CRITICAL] MSSQL Dockerfile Missing ODBC Driver

apps/api/Dockerfile builds from python:3.12-slim and installs Python deps, but never installs the Microsoft ODBC Driver 17/18 for SQL Server. The aioodbc package requires this system library at runtime. Production containers will crash on startup with 'ODBC Driver not found'. FIX: Add RUN apt-get install -y unixodbc-dev + curl/gpg steps to install msodbcsql18.

[CRITICAL] In-Process Worker Has No Restart Mechanism

In dev mode (QUEUE_PROVIDER=memory), the worker coroutine is launched once in the FastAPI lifespan hook. If it raises an unhandled exception, it stops silently. New scan requests are enqueued (HTTP 202 returned) but never processed. The API gives no indication of this state. FIX: Wrap the worker loop in a restart supervisor, OR add a /health check that reports worker liveness, OR document clearly that dev mode is not reliable for extended use.

[CRITICAL] No CI/CD Pipeline

Zero GitHub Actions workflows or Azure Pipelines found. Every deployment is manual. This means broken code can reach production undetected and there is no automated test gate. FIX: Add a basic CI pipeline: lint -> unit tests -> build Docker image -> push to registry.

5.2 High Severity Issues

[HIGH] No API Rate Limiting

The FastAPI backend has no rate limiting on any endpoint. A malicious or buggy client could flood /scans/run, causing runaway OpenAI token usage and Graph API throttling. FIX: Add slowapi (FastAPI rate limiter) or an API Gateway layer with per-user limits.

[HIGH] MCP Multi-User Auto-Resolution is Ambiguous

The MCP server auto-resolves user_id by querying the DB for 'the first user' when no explicit user_id is provided. In a multi-tenant deployment this is dangerous - the wrong user's tasks could be returned or modified. FIX: Require explicit user_id/tenant_id in all MCP tool calls, or tie MCP sessions to authenticated users via the X-Api-Key header.

[HIGH] PKCE Not Enforced in OAuth Flow

The OAuth flow uses MSAL's confidential client. While a client secret is present, PKCE (code_challenge/code_verifier) is not explicitly enforced. For public-facing deployments this is a defense-in-depth gap. FIX: Enable PKCE in MSAL config even for confidential clients.

[HIGH] PII in Application Logs

User email addresses, task titles (potentially sensitive content), and sender names are logged at INFO/DEBUG level throughout the scan pipeline. If logs are shipped to a SIEM or log aggregator, this constitutes uncontrolled PII exposure. FIX: Scrub or hash PII fields before logging. Log user_id/tenant_id instead of email.

5.3 Medium Severity Issues

[MEDIUM] No Optimistic Locking on Task Creation

Concurrent scan triggers (manual + scheduled) can create duplicate tasks if both pass the dedup check before either has committed. SQLAlchemy's flush() helps but doesn't guarantee isolation at the database level without a unique constraint. FIX: Add a UNIQUE constraint on (tenant_id, user_id, source_message_id) in the tasks table.

[MEDIUM] Excel Sync Has No Conflict Resolution

The Excel sync appends rows to TaskInbox.xlsx. If the user edits the workbook (changes a task status, adds notes), the next sync will append a duplicate row with the original data. There is no read-back and reconcile step. FIX: Before appending, read existing row IDs from the table and skip already-synced tasks, OR switch to update-by-row-ID approach.

[MEDIUM] APScheduler TZ Handling May Drift

The scheduler fires once per minute and checks each user's preferred daily scan time in their local timezone. Timezone data depends on the pytz/zoneinfo db in the container. DST transitions can cause double-fires or skips. FIX: Use an idempotent 'last_ran_date' check so re-firing on the same day is harmless.

[MEDIUM] Frontend Has No Error Boundaries

Next.js app has a single /auth/error page but no React error boundaries in the component tree. An unhandled promise rejection in any SWR fetcher crashes the entire page to a blank screen with no user feedback. FIX: Wrap route layouts in error.tsx files (Next.js 14 convention) with user-friendly fallback UI.

[MEDIUM] No Pagination on /tasks Endpoint

GET /tasks accepts limit/offset params but the default limit is not validated server-side. A request with limit=100000 could return the entire tasks table, causing memory spikes. FIX: Cap limit at 500 server-side and document the maximum in OpenAPI schema.

[MEDIUM] Attachment URL Orphaning

When a message is scanned, email attachments are archived to local/blob storage. If the user later deletes the original email, the attachment URL in the DB becomes an orphan. If local storage is used and the container is recreated, all local attachments are lost. FIX: Always use Blob storage for attachments, even in dev; add a health-check for orphaned URLs.

6. Performance Analysis

Latency, throughput, and scalability concerns

6.1 Graph API Scanning Efficiency

Each scan issues sequential paginated GET requests to /me/messages (50 per page). For a user with 500 new messages, this is 10 sequential HTTP round-trips before AI extraction even begins. Graph supports $batch requests (up to 20 ops per batch) and delta tokens for incremental sync - neither is used.

- CURRENT: Sequential paged GET requests - O(N) round-trips where N = message count / 50.
- BETTER: Graph $batch to parallelize page fetches.
- BEST: Graph delta tokens (Change Tracking) - only fetch truly new messages since last sync token.
- AI extraction is sequential per message (one OpenAI call per message). For 100 messages this could take 2-3 minutes with default timeouts.
- IMPROVEMENT: Batch messages into groups and use asyncio.gather() for parallel AI calls (limited to e.g. 5 concurrent to avoid OpenAI rate limits).

6.2 Database Query Patterns

All queries filter by (tenant_id, user_id). Without composite indexes on these columns, every query does a full table scan as data grows.

- Missing index: (tenant_id, user_id) on tasks, source_messages, scan_runs, audit_logs.
- Missing index: (tenant_id, user_id, status) on tasks - used in /tasks/today and /tasks/overdue.
- Missing index: body_hash on source_messages - used in dedup layer 3.
- The Alembic migration creates tables but defines no explicit indexes beyond primary keys.
- FIX: Add explicit index definitions in the Alembic migration or models.py Index() declarations.

6.3 Frontend Data Fetching

SWR is used correctly for data fetching. However, the dashboard page fetches 5+ separate endpoints on load (tasks, scans, settings, connections, stats) with no parallel aggregation endpoint. Each SWR hook fires independently.

- 5 sequential render-triggered fetches create a waterfall effect on initial page load.
- SWR's refreshInterval is not set - data staleness after a scan completes requires manual refresh.
- IMPROVEMENT: Add a /api/dashboard/summary endpoint that aggregates common dashboard data in one DB query.
- IMPROVEMENT: Set SWR refreshInterval=30000 on scan status hooks to auto-update after scan completion.

6.4 Token Usage & Cost Control

Every scanned message that passes noise/dedup filters makes one OpenAI API call. There is no caching, batching, or cost cap.

- No token count logging - impossible to audit OpenAI costs per user or per scan.
- No per-user or per-day token budget cap.
- Re-scanning after a config change will re-charge tokens for already-processed messages (dedup prevents duplicate tasks but doesn't prevent duplicate AI calls if source_message records are cleared).
- IMPROVEMENT: Log prompt_tokens + completion_tokens per scan_run in the scan_runs table.
- IMPROVEMENT: Add a setting for max_messages_per_scan to cap cost per run.

7. Security Review

Authentication, authorization, data protection, and threat surface

7.1 Strengths

- OAuth 2.0 / MSAL confidential client - correct delegated-auth pattern for Microsoft 365.
- Fernet encryption on stored Graph tokens - tokens at rest are encrypted with a key that is separate from the database.
- httpOnly JWT session cookie (mtr_session) - XSS cannot steal the session token via JavaScript.
- JWT with 8-hour expiry - reasonable session lifetime.
- Tenant + user ID filtering on every query - strong multi-tenant isolation.
- Audit log table captures all user actions with timestamps.
- HTML body stripped before AI prompt - prevents HTML-injection in prompts.

7.2 Vulnerabilities & Gaps

[HIGH] SECRET_KEY / JWT_SECRET Default Values in .env.example

The .env.example ships with SECRET_KEY=dev-secret-change-me and JWT_SECRET=dev-jwt-change-me. If a developer copies .env.example to .env without changing these, production JWTs are signed with a known key. Anyone can forge session tokens. FIX: Generate secrets programmatically at first run, or fail startup if defaults are detected in non-dev environments.

[HIGH] CORS Configuration Not Visible

FastAPI's CORSMiddleware configuration is not explicitly shown. If CORS is set to allow all origins (allow_origins=['*']) with allow_credentials=True, this violates the CORS spec (browsers reject credentialed requests to wildcard origins) and could allow cross-origin session hijacking if misconfigured. FIX: Ensure CORS is restricted to FRONTEND_URL. Validate this is enforced.

[MEDIUM] TOKEN_ENCRYPTION_KEY Auto-Generated in Dev

When TOKEN_ENCRYPTION_KEY is absent, the app auto-generates a Fernet key at startup. On container restart, the key changes and all stored tokens become unreadable. This silently breaks all existing users' Graph connections without an actionable error. FIX: Require TOKEN_ENCRYPTION_KEY to be set explicitly in all environments, or persist the auto-generated key to a file/Key Vault.

[MEDIUM] MCP API Key is Static

The MCP server is protected by a single static X-Api-Key header (MCP_API_KEY env var). There is no per-client key rotation, no key expiry, and no audit of which client called which tool. A leaked MCP key grants full access to all tasks across all users. FIX: Issue per-client API keys stored in DB, add key rotation endpoint, log caller identity in audit logs.

[MEDIUM] No CSRF Protection on State-Changing Endpoints

POST/PATCH/DELETE endpoints rely on the JWT cookie for auth. Without CSRF tokens or SameSite=Strict cookie policy, a malicious page could trigger state-changing requests from a victim's browser. FIX: Set cookie SameSite=Lax or Strict, OR add CSRF token validation on state-changing endpoints.

[LOW] JWT Algorithm HS256 with Shared Secret

HS256 (HMAC-SHA256) with a shared secret means anyone with the secret can both verify AND forge tokens. RS256 (asymmetric) would limit token forgery even if a verification key leaks. This is acceptable for an MVP but should be upgraded for enterprise deployments.

[LOW] No Input Sanitization on Task Search

GET /tasks/search?q=... passes the query string to a SQL LIKE clause. SQLAlchemy parameterizes this correctly (no SQL injection risk), but there is no length cap on the query parameter. A very long query string could cause unnecessary DB load. FIX: Cap search query length at 200 characters server-side.

8. What's Missing & Lagging Behind

Incomplete features, stubs, and Phase 2 items

8.1 Incomplete Implementations (Stubs)

Teams Scanning	services/graph/teams.py exists and has the Graph API calls, but enable_teams_scan defaults to false. The UI has a Teams settings page but it is read-only. No end-to-end test covers Teams scanning. This is the most impactful missing feature for a Microsoft 365 product.
Realtime Webhooks	services/graph/webhooks.py is a near-empty stub. Graph subscriptions (for push notifications on new email) are not configured, validated, or handled. Users must manually trigger scans or rely on the daily scheduler.
Semantic Deduplication	Only hash-based dedup is implemented. Embedding-based similarity (e.g., using text-embedding-3-small) would catch rephrased duplicates. Architecture is ready; implementation is missing.
Analytics / Insights UI	No chart or trend data is shown in the UI. The scan_runs table has metrics (messages_scanned, tasks_created) but they are not visualized anywhere.
Jira / ClickUp Connectors	The sync architecture is generic (pluggable targets) but only Excel and Planner adapters are implemented. Third-party task manager integrations are absent.
Mobile / Responsive UI	The Next.js frontend uses Tailwind but no responsive breakpoints are explicitly tested. The sidebar layout likely collapses poorly on small screens.
Email Digest / Notifications	No outbound notifications (email digest, Teams message, push) when new tasks are discovered. Users must actively open the app.
Admin Panel	No tenant-admin view. An admin cannot see all users, manage quotas, or revoke connections without direct DB access.

8.2 Missing Infrastructure

- CI/CD Pipeline: No automated test-on-PR, no Docker image build, no deployment automation.
- Database Indexes: No performance indexes defined in migrations (see Section 6.2).
- Health Probes in Docker Compose: No healthcheck: directives on the api or worker services.
- Graceful Shutdown: No SIGTERM handler to drain the in-process queue before shutdown.
- Log Aggregation: Logs go to stdout/stderr. No structured JSON logging or aggregation pipeline.
- Secret Rotation: No mechanism to rotate TOKEN_ENCRYPTION_KEY or JWT_SECRET without downtime.
- Backup Policy: No automated DB backup for MSSQL in Azure. SQLite dev DB is ephemeral.
- Load Testing: No load or stress test suite to validate throughput before production launch.

8.3 Frontend Gaps

- No loading skeletons - pages show blank white during SWR fetch (poor perceived performance).
- No empty-state illustrations - empty task list shows plain text, not a friendly onboarding prompt.
- No keyboard navigation / accessibility audit (ARIA labels, focus management).
- No internationalization (i18n) - all strings are hardcoded English.
- No dark mode (Tailwind supports it trivially but it is not wired up).
- Task detail view is read-only in some fields - users cannot edit due date or reassign.
- No bulk actions - user cannot select multiple tasks and archive/approve them at once.
- Settings pages for Teams, Excel, Planner have no connection-test or validate button.

9. Recommendations & Improvement Roadmap

Prioritized action plan - immediate, short-term, and strategic

Immediate (Pre-Production Must-Fixes)

1. Fix Dockerfile: Install Microsoft ODBC Driver 18 for SQL Server in the production image.
2. Add CI/CD: Create a GitHub Actions workflow with: lint -> pytest -> docker build -> push.
3. Enforce strong secrets: Fail startup if SECRET_KEY or JWT_SECRET equals the example default value in non-dev environments.
4. Add API rate limiting: Integrate slowapi with per-user limits on /scans/run (e.g., 5/min) and /tasks/search (e.g., 30/min).
5. Add DB indexes: Add composite indexes on (tenant_id, user_id) + status + body_hash via a new Alembic migration.
6. Fix worker liveness: Add /health?deep=true that reports worker coroutine alive/dead status.
7. Fix MCP user resolution: Remove the 'first user in DB' fallback; require explicit auth.
8. Set cookie SameSite=Lax: Ensure the mtr_session cookie has SameSite=Lax at minimum.

Short-Term (Sprint 1-2)

1. Implement Teams scanning: Enable end-to-end (Graph calls already written), add tests, expose enable_teams_scan as a per-user setting in the UI.
2. Add Graph delta tokens: Replace high-water-mark polling with delta token change tracking for Outlook - dramatically reduces Graph API calls.
3. Parallel AI extraction: Use asyncio.gather() with a semaphore (max 5 concurrent) to parallelize OpenAI calls per scan run.
4. Log token costs: Capture prompt_tokens + completion_tokens in scan_runs table; display a 'tokens used this month' stat on the settings page.
5. Add Next.js error.tsx boundaries: Wrap each route with an error boundary showing a user-friendly message instead of a blank screen.
6. Add loading skeletons to dashboard and task list pages.
7. Add unique constraint: (tenant_id, user_id, source_message_id) on tasks table to prevent concurrent-scan duplicates at the DB level.
8. Scrub PII from logs: Replace user emails with user_id in log statements throughout the scan pipeline.

Strategic (Phase 2)

1. Realtime webhooks: Implement Graph subscriptions so new emails trigger immediate extraction rather than waiting for the next scheduled scan.
2. Semantic deduplication: Integrate text-embedding-3-small to compute vector similarity for near-duplicate detection across rephrased messages.
3. Analytics dashboard: Visualize scan history, tasks-per-day trend, resolution rate, and top senders using the existing scan_runs metrics.
4. Admin panel: Build a tenant-admin view for user management, quota enforcement, and connection health monitoring.
5. Jira / ClickUp connectors: Extend the pluggable sync architecture to third-party task managers.
6. Per-user noise rules: Allow users to whitelist/blacklist senders and subject patterns from the settings UI.
7. Upgrade JWT to RS256: Issue signed tokens with asymmetric keys for stronger forgery resistance.
8. Per-client MCP keys: Replace the single static API key with per-integration keys stored in DB with rotation and audit logging.
9. Mobile-responsive UI: Add Tailwind responsive breakpoints and test on mobile viewports.
10. SLA / cost guardrails: Add per-user daily token budget and scan frequency caps to control OpenAI costs at scale.

10. Test Coverage Analysis

What is tested, what is not, and what needs real credentials

10.1 Coverage Summary

Total test files	15
Estimated tests	37+
Test framework	pytest + pytest-asyncio (asyncio_mode=auto)
HTTP mocking	respx 0.21.1 (HTTPX transport-level mocking)
Time mocking	freezegun 1.5.1
DB backend	aiosqlite in-memory (no external DB needed for tests)
OpenAI mocking	unittest.mock.AsyncMock on AsyncAzureOpenAI.chat.completions.create
Graph mocking	respx route matching on Graph API base URL

10.2 Well-Tested Areas

- OAuth callback flow: state validation, code exchange, tenant upsert, user creation.
- JWT cookie issuance and validation (test_auth_cookie.py).
- Deduplication: all 3 layers tested with collision scenarios (test_dedup.py).
- Noise filtering: newsletter, OOO, marketing patterns (test_noise.py).
- AI extraction schema validation: valid/invalid JSON shapes, confidence thresholds (test_extractor_validation.py).
- End-to-end scan flow: 6 scenarios including FYI email, low confidence, attachment handling, dedup collision, and AI failure (test_scan_flow_integration.py).
- Tenant scoping: verifies cross-tenant data isolation (test_tenant_scoping.py).
- MCP tools: all 9 tools tested with mocked DB responses (test_mcp_tools.py).
- Health + readiness endpoints (test_health_mcp_http.py).

10.3 Missing / Weak Test Coverage

- Teams scanning: zero test coverage (feature disabled, but at least a mock test should exist).
- Excel sync: no test for the append-rows path or conflict scenarios.
- Planner sync: no test for approval gate, plan discovery, or task creation.
- Token refresh: no test for the 60-second pre-expiry refresh path.
- Queue adapter: no test that Service Bus adapter correctly serializes/deserializes scan jobs.
- Scheduler: no test that the daily trigger fires at the correct user timezone.
- Noise filter false-positive scenarios: no test for a legitimate task email that contains newsletter-like keywords (e.g., subject contains 'update').
- Rate limiting (once added): no tests for 429 responses.
- Frontend: zero automated tests (no Jest, no Playwright, no Cypress).
- Performance / load: no benchmark or stress test.

10.4 Tests Requiring Real Credentials

The following flows are fully mocked in CI but must be validated manually with real Azure credentials before go-live:

- Microsoft Entra ID OAuth - real Azure AD app registration, tenant consent, token exchange.
- Outlook scanning - real Exchange Online mailbox with delegated Mail.Read permission.
- Teams scanning - real Teams tenant with Channel.ReadBasic.All + Chat.Read permission.
- Azure OpenAI GPT-5.2 - real endpoint, deployment name, and API key.
- Excel workbook - real OneDrive with Files.ReadWrite, table creation and row append.
- Planner - real Microsoft Planner plan/bucket with Tasks.ReadWrite permission.
- Azure Blob Storage - real container with write access for attachment archival.
- Azure Service Bus - real queue with send/receive permissions for worker handoff.

11. Final Summary & Conclusion

Mela Task Radar is an impressively complete MVP for an AI-powered task intelligence platform. The architecture is clean, async-first, and well-layered. The core scan pipeline - from Graph API ingestion through noise filtering, deduplication, AI extraction, and multi-target sync - is implemented end-to-end and covered by a solid test suite.

The project is genuinely close to production-ready. However, several blockers must be addressed before it can safely serve real users:

Stop-Ship Items

- ODBC driver missing from Dockerfile -> production containers crash.
- In-process worker has no liveness/restart -> scans silently die.
- No CI/CD -> broken code can reach production undetected.
- Default secret values -> JWT forgery risk if not changed.
- No API rate limiting -> runaway OpenAI costs and Graph throttling.

High-Value Quick Wins

- Graph delta tokens -> 10x reduction in scanning latency and API quota usage.
- Parallel AI extraction -> 5x faster scan runs for large mailboxes.
- DB indexes -> prevents query degradation as task count grows.
- SWR refresh intervals -> dashboard auto-updates after scan completes.
- Loading skeletons -> dramatically improves perceived frontend performance.

Strategic Differentiators to Build Next

- Teams scanning - the single most impactful feature gap for a Microsoft 365 product.
- Realtime webhooks - transforms the product from 'daily digest' to 'live task radar'.
- Analytics dashboard - makes the value visible (how many tasks found, resolution rate).
- Per-user noise rules - reduces false positives, a key UX pain point.

Closing Note

The engineering quality is high for an MVP. The codebase is readable, well-typed, and consistently structured. Documentation is comprehensive and accurate. With the stop-ship items resolved and real-credential validation completed, this product is ready for a limited pilot launch. Phase 2 (Teams, webhooks, analytics) would transform it from a useful tool into a genuinely compelling product.

Total estimated effort to reach production-safe state: 3-5 engineering days.